

## Conceptual Database Design

### Entity/Relationship (E/R) Model

1

## Database Design


- Requirements collection and analysis
  - Determine what information the database must hold
  - Determine the relationships among the components of that information
- *Conceptual database design* using some *data model*
- A **data model** is a collection of concepts for describing:
  1. data and relationships among data
  2. data semantics and constraints

2

## Design Approaches & Notations

- Entity-Relationship Model (E/R)
- Object-Oriented (e.g., Object Definition Language -- ODL)
- Semi-structured data (e.g., XML)

### A Design Process



3


## Entity-Relationship (E/R) Model

- E/R model (Chen 1976) is a *graphical approach to database modeling*
- It grew out of modeling applications
- It is widely used in database design
- There is no one standard for the E/R language/notation

4

## Entity-Relationship (E/R) Model

- Collection of abstraction / modeling primitives



5


## Entities and entity sets

- **Entity** -- A real world object that is "distinguishable" from other objects
- **Entity Set** -- A collection of similar entities
  - All entities in an entity set have the same set of **attributes**
- In ODL:
  - **Object** corresponds to **entity**
  - **Class** corresponds to **entity set**



6


## Entities and entity sets



7


## Relationships and Relationship Sets

- Relationships are association among entities
- Relationship set is a set of relationships of the same type
  - If  $E_1, \dots, E_n$  are entity sets, a relationship  $R$  on these sets is defined as:  $R \subseteq E_1 \times \dots \times E_n$ 
    - In general, relationships are n-ary, where  $n \geq 2$
    - Many database relationships are binary



8


## Relationships and Relationship Sets



9

## Multiplicity of Binary Relationships


- EnrolledIn relationship between entity sets Student and Course is many-many



10

## Multiplicity of Binary Relationships


- In E/R diagrams, arrows can be used to indicate the multiplicity of a relationship



- If a relationship is many-one from entity set Course to Department, we place an arrow entering Department
- The arrow indicates that each entity in the entity set Course is related to at most one entity in the entity set Department
- In this case, we may also say that the relationship from Department to Course is one-many (also shown as 1-M).

11

## Multiplicity of Binary Relationships




- The arrow indicates that each entity/tuple in relation Movie is "related to" at most one entity in Studio
- Owns relationship is one-many from Studio to Movie

12


## Multiplicity of Binary Relationships

- The relationship **Chair** between the entity sets **Department** and **Professor** is **one-one**, indicating which professor is the chair of which department. Note that this represents also the situation where a department is assigned no chair.



13

## Multiplicity of Binary Relationships




14

## Multiway Relationships

- ODL** (an OO notation which we will introduce later) allows defining only **binary relationships**, i.e., relationships involving two classes.
- In general, we need to represent **n-ary** (multi-way) relationships, i.e., relationships involving more than two entity sets
- E/R** model allows defining n-ary relationships conveniently.
- An n-ary relationship in **E/R** is represented by a diamond connecting all entity sets involved.

15


## Multiway Relationships



16

## Attributes on Relationships


- Consider **EnrolledIn** relationship between **Student** and **Course**



- We may wish to record student's **grade** for this course
- Where should it go?
- Sometimes, it is more appropriate to associate **attributes** with a **relationship** rather than with the entity sets involved

17

## Attributes on Relationships



- Suppose we want to record the star's **salary** associated with this contract
- Where should it go?

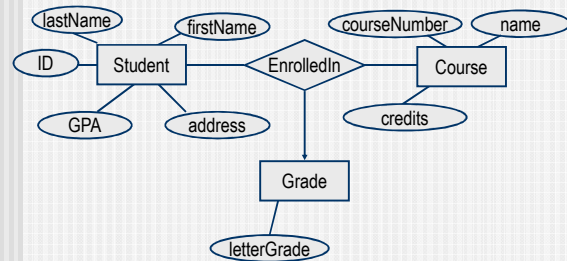
18

## Moving an attribute to an entity set

- We could **add an attribute(s) to a relationship**
- Alternatively, we could do the following:
  - invent a **new** entity set, whose entities have the attributes ascribed to the relationship
  - "connect/include" this entity set in the relationship
  - omit the attribute from the relationship itself
- Consider again the "salary" for the "Contract" example:

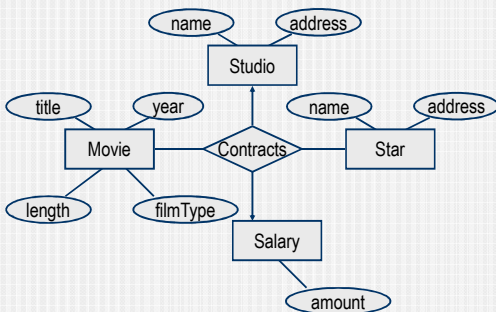
19

## Moving an attribute to an entity set



20

## Moving an attribute to an entity set




21

## Roles in Relationships

- It is possible that one entity set appears **two or more** times in a relationship
- Suppose, we want to capture relationship between two **movies**, one of which is the **sequel** of the other

22

## Roles in Relationships




Each line to the entity set represents a different **role** that the entity set plays in the relationship

23

## Roles in Relationships

Each line to the entity set represents a different **role** that the entity set plays in the relationship




24



## A more complex example

- Suppose, each star is associated with exactly one studio
- Supp. studio s1 of star a1 may further contract with another studio s2 to allow a1 to play in movie m2 made by studio s2



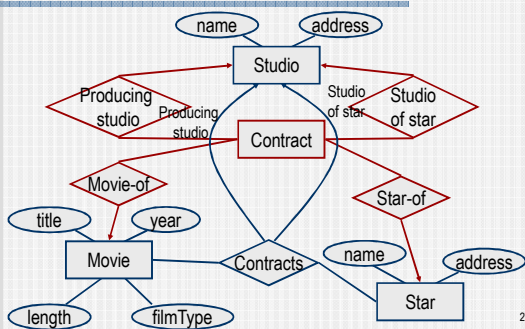
25

## Converting n-ary relationship to binary

- Any n-ary relationship  $R$  can be converted into a collection of  $M-1$  binary relationships without losing the information represented by  $R$
- To do this:
  - Introduce a new entity set  $E$ , called **connecting** entity set, whose entities might be thought of as tuples in  $R$  (the original n-ary relationship)
  - Introduce  $M-1$  relationships from  $E$  to each one of the entity sets involved in  $R$
  - If an entity set  $E$  plays **more than one** role, then  $E$  is the target of one relationship for **each** role

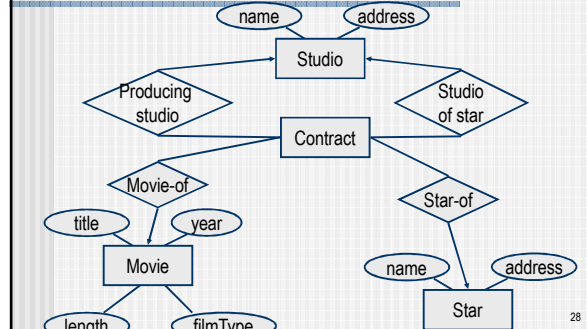
26

## Converting n-ary relationship to binary



27

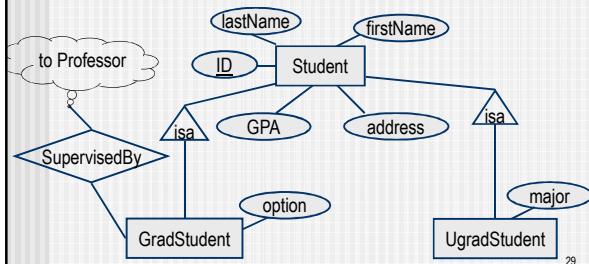
## Converting n-ary relationship to binary



28

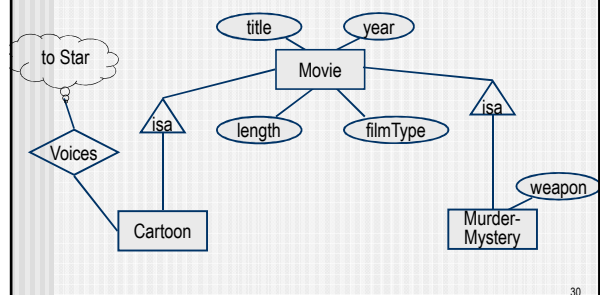
## Inheritance in E/R

- Inheritance in E/R is expressed by **isa** relationship



29

## Inheritance in E/R



30

## Inheritance in E/R

- There is a subtle difference between the concept of inheritance in an OO approach and in E/R
- In OO, an object must be a member of exactly one class
- In E/R
  - We consider an entity as having "components" belonging to several entity sets that are "part of" a single **isa**-hierarchy
  - The "components" are connected into a single entity by the **isa** relationships
  - The entity has whatever attributes any of its components has, and participates in whatever relationships its components participate in
  - In an E/R diagram, we represent an entity set (e.g., **CartoonMurderMystery**) only if it has attributes and/or relationships of its own

31

## Constraints

- There are some important aspects of the real world that cannot be represented using the ODL or E/R model introduced so far
- The additional information about these aspects often takes the form of constraints on the data
- Sometimes modeling this additional information goes beyond the structural and type constraints imposed by classes, entity sets, attributes, and relationships


32

## A Classification of Constraints

- A **Key K** is a set of attribute(s) that uniquely identifies an object within its class or an entity within its entity set R; Defn:  $K \subseteq R$ .
  - That is, no two entities may agree in all their key values
- Single-value constraints** are requirements that the value in a certain context/role be unique. In addition to key constraints, there are other sources of single-value constraints (e.g., M:1 or 1:1 relationships, like a department's chair)
- Referential integrity constraints** are requirements that a value referred to by some object/entity must actually exist in the database; This means, no dangling pointers
- Domain constraints** require that the value of an attribute must be drawn from a specific set of values (called attribute domain), or lies within a specific range
- General constraints** – arbitrary assertions that must hold on the db

## Keys


- A **superkey** is a set of attributes whose values uniquely identify an entity (object) in the entity set (class); this set may not be minimal.
- A **minimal superkey** is called a (candidate) key.
- An entity set may have more than one candidate key. One of them is picked (by ?) as the **primary key**; others may be called **alternates**
- In E/R, we underline the attributes forming a key of an entity set
- No notation in E/R for alternate keys



Here, ID is the key of **Student**

## Example

- What should we consider as a key for **Movie** ?
- title**?
  - there could be different movies with the same name
- {title, year}**?
  - there still could be two movies with the same title made in the same year, but that's very unlikely as we understand!



35

## Example

- What should be the key for **Star = {name, address}**?
- name**?
  - We may think that the name cannot serve to distinguish two people, but for stars, the names distinguish them, since traditionally they choose "stage names" as names



36


## Example

- What should be considered as the key for **Studio**?
- name**?
  - It is "reasonable" to assume that there are **no** studios with the **same names**



37

## Selecting Primary Key



- Suppose candidate keys for E are :
- {A, B}
  - {D, N}
  - {G, M, N}

- Which of the three should be pick as the **primary key**?



38

## Selecting Primary Key

- Some criteria to choose a **primary key** when there are **alternate keys** (i.e., more than one candidate):
  - Total size of the attributes forming a key
  - Number of attributes forming a key
  - Convenience/Natural choice
  - A combination of the above

39

## Single-Value Constraints

- In E/R:
  - attributes are **atomic** (first normal form, 1NF)
  - an **arrow (→)** can be used to express the multiplicity
  - What about multi-valued or structured in E/R?
    - No for attributes but Yes for relationships

40


## Single-Value Constraints

- Suppose A is a single-valued attribute.
- Using the E/R model introduced so far, we can't:
  - Express that the value for A **must** be present (e.g., when A is the key or part of a key), or
  - Express that the value for A **may** be present (e.g., when A's value is optional, for which we use NULL)
- If the choice is not explicit, we use the following "defaults":
  - The value for A must exist if A is part of the key
  - The value for A is optional, otherwise

41

## Referential Integrity Constraints


- For relationships:
  - Single-value + Existence = Referential Integrity Constraint**
  - We **extend the arrow notation** to indicate a reference is **mandatory** (to support referential integrity)



- This means, there is **no course listed in the database unless there is a department giving the course.**

42

## Referential Integrity Constraints



- The studio for a movie must always be present in **Studio**
- If someone is a president, then the studio that he/she runs must exist in **Studio**
- However, the model allows studios without presidents (temporarily)

43


## Domain constraints

- **Domain constraints** restrict the values of an attribute to be drawn from a set
- E/R does **not** support imposing domain constraints
- ODL allows using **types** to limit/control possible values of the attributes
- ODL does not support restricting the "range" of values allowed for an attribute

44

## Relationship degree constraints


- **Relationship degree constraints** (multiplicity) restrict the number of entities in the entity sets involved in a relationship
- For example, we can impose a constraint that:  
*a student cannot be enrolled in more than 5 courses*
  - In E/R, we may attach a "bounding number" to the corresponding link
  - In ODL, a set of references or an **array** of size 5 (of reference type) will do



45

## Weak Entity / Relationship Sets: Example

- Log records the transactions initiated through an ATM
- Each transaction has a number, a date, and an amount
- Different accounts might have transactions by the same number, on the same date, and for the same amount



46

## Weak Entity / Relationship Sets

- A **strong** entity set has a **key**
- A **weak** entity set does not have sufficient attributes of its own to form a **key**. It participates in a M-1 relationship (with no descriptive attributes) with a strong entity set
- **Discriminator** of a weak entity set is a set of attributes that distinguishes among the entities corresponding to a strong entity
- **key** of a weak entity set = key of the strong entity together with the discriminator of the weak entity
- In E/R, these are represented by:



47

## Another example

- A movie studio may have several film crews
- The crews in a studio may be designated as crew 1, crew 2, and so on
- All studios use the same designations for crews




48



## Sources of Weak Entity Sets

- Entity sets fall into a **hierarchy**
  - The entities in set **Crew** are **sub-units** of entities in set Studio
    - The "numbers" used for crew entities are not unique until we take into account the name of Studio with which she/he is associated (subordinate)
  - It is possible to have a chain of "weak" entity sets/relationships



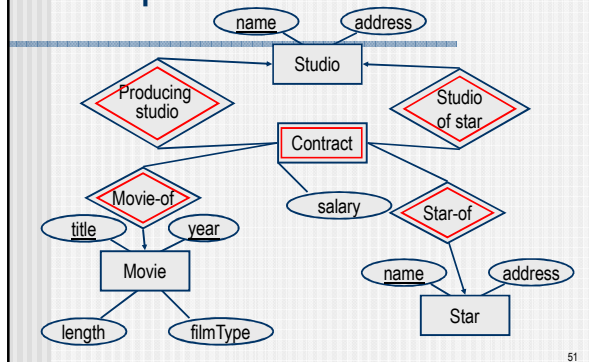
49

## Sources of Weak Entity Sets

- Connecting entity sets (CES):**
  - CES's often have no attributes of their own
  - Their key is formed from the key attributes of the entity sets which they connect
- A connecting entity set is always weak

50

## Example



51

## Design Principles

- Design should
  - Reflect the "reality" we are trying to model – **faithful/realistic**
  - Avoid redundancy -- **minimal**
    - Redundant information takes space
    - Could cause inconsistency
  - Be as **simple** as possible
- Be careful when choosing between using attributes and using classes or entity sets. Remember that:
  - An attribute is **simpler** to implement than a class/entity set or a relationship
  - If something has more information associated with it than just its name, it should probably be modeled as an entity set or a class

## A quick test!

- Which of the following statements is **NOT** correct?
  - A. A data model is a collection of concepts for describing data and their relationships.
  - B. A data model is a language for describing the data semantics and the constraints.
  - C. A data model is a language for expressing queries and transactions over a database.**
  - D. A data model is an abstract representation of the structure of the data.

53